

The Equilibrium Is the Initialization: Lazy Identity Collapse in Physics-Structured Deep Equilibrium Reasoning

Joyjeet Singh
Independent Researcher
ORCID: 0009-0005-1512-7439

July 13, 2026

Abstract

Deep equilibrium models promise input-adaptive implicit computation: harder problems should demand more solver iterations, and the solved equilibrium should encode the result of genuine iterative inference. We report a cautionary study of a port-Hamiltonian DEQ with a learned initialization on two reasoning tasks — ProofWriter entailment over frozen DeBERTa embeddings and a BFS-verified graph-reachability benchmark — in which the implicit computation is a silent no-op. Across tasks, seeds, and controlled ablation arms, the solved equilibrium equals the solver’s *start point* to numerical precision, and bypassing the solver entirely changes test accuracy by +0.00 percentage points in 18 of 19 training runs. Controlled interventions falsify the tempting explanation: removing the anchoring term reproduces every result, and retraining with noise-decoupled starts yields a solver that converges to the noisy start while the decoder learns to ignore it. The single escaping run diverges instead ($\|h^* - z_0\| = 171$), producing a co-adapted noise channel whose removal *improves* accuracy. Iteration counts are uncorrelated with ground-truth difficulty ($r = 0.009$), and the full apparatus never outperforms a two-layer MLP on either task. We trace the mechanism to gradient starvation along two distinct routes, show that the standard zeroing ablation is confounded and gives wildly seed-dependent answers where the correct substitution test gives a stable zero, and distill a four-test diagnostic protocol for auditing claimed implicit computation. All experiments run on a single free Colab GPU; code, raw logs, and analysis scripts are released.

1 Introduction

Implicit and recurrent depth is one of the field’s recurring hopes for reasoning: rather than fixing computation at a static layer count, let the model iterate — a deep equilibrium model (DEQ) solving for a fixed point [1], a looped or recurrent-depth transformer unrolling until done [2, 3] — and let harder inputs consume more computation [4]. Physics-structured parameterizations make the proposal more attractive still: port-Hamiltonian and symplectic constraints promise stable, energy-controlled latent dynamics on which an implicit solve can safely iterate [5–7]. The promise is concrete and checkable: the solved equilibrium should carry the *result* of iterative inference, and solver effort should track problem difficulty.

This paper documents, in controlled detail, a design in which every part of that promise fails silently. We study a port-Hamiltonian DEQ whose state is initialized by a learned, goal-conditioned feed-forward network — a natural pattern whenever an implicit layer is grafted onto a pretrained encoder, and one that (with or without an explicit anchoring term pulling the solve toward the

initialization) appears in many variants across the implicit-models literature. Training converges, accuracies look respectable, and solver telemetry looks plausible. Yet the equilibrium is a copy of its initialization: across two tasks, nineteen training runs, an anchor-strength sweep, and three intervention arms, the fixed-point solve moves the state by at most 10^{-6} and contributes exactly +0.00 percentage points of accuracy, while iteration counts ignore a ground-truth difficulty signal entirely.

The obvious mechanistic story — the anchor both starts and attracts the solver, so the dynamics never receive off-anchor gradients — turns out to be wrong, and we consider its falsification the most useful part of the study. Removing the anchor changes nothing; retraining with noise-decoupled starts produces a solver that faithfully converges to the *noisy* start while the decoder learns to ignore the channel; and the one run that escapes the identity regime does so by diverging, yielding a co-adapted noise channel whose removal improves accuracy. The correct diagnosis is gradient starvation: both tasks are solvable through the direct context pathway — a two-layer MLP baseline matches or beats the full apparatus throughout — so “compute nothing” is a stable optimum, reached by two different gradient routes in our two pipelines. Along the way we show that the standard zeroing ablation is confounded for implicit modules, producing seed-dependent deltas as large as +44pp where the correct substitution test gives a stable zero.

Our contributions: (i) a mechanistically dissected case study of silent no-op implicit computation, with the anchor-causation hypothesis explicitly tested and falsified by controlled retraining arms; (ii) identification of two degenerate regimes — lazy identity collapse and solver blow-up — with per-seed evidence that pooled statistics would obscure; (iii) a demonstration that zeroing ablations confound pathway with computation, and the substitution test that deconfounds them; (iv) a four-test diagnostic protocol (Section 7) cheap enough to run on a free GPU and applicable to DEQs, looped, and recurrent-depth models generally; and (v) a fully reproducible negative result: every number traces to released raw logs via a released analysis script.

2 Related Work

Deep equilibrium models and implicit layers. Deep equilibrium models compute their output as the fixed point of a learned layer, differentiated implicitly [1, 8, 9], with guarantees available under monotonicity assumptions [10] and practical solves typically using Anderson acceleration [11]. Known DEQ pathologies concentrate on the *unstable* side: growing iteration counts and brittle convergence, addressed by Jacobian regularization [12], by cheaper approximate gradients [13], and by encouraging path independence, which correlates with the ability to exploit extra test-time iterations [14]. Our study documents the opposite, quieter failure: *degenerate convergence*, in which the solve succeeds trivially because the equilibrium is the initialization. Path independence offers a useful lens here — a model whose output ignores initialization because it computes something is the desirable extreme; ours ignores computation because it copies the initialization.

Physics-structured latent dynamics. Encoding mechanics into network dynamics — Hamiltonian neural networks [5], Lagrangian variants [15], symplectic parameterizations [6], and port-Hamiltonian extensions with explicit dissipation [7], in the broader family of continuous-depth models [16] — promises stable, energy-controlled latent evolution. Our dynamics use the standard structure–dissipation decomposition with learned kinetic and potential energies and a context-conditioned damping term. The irony our results add: the same properties that make such dynamics stable to integrate make “stay where you started” an easily reachable, easily maintained solution when nothing in the loss demands motion.

Recurrent depth, looped models, and test-time compute. Allocating more sequential computation to harder inputs is an old and newly urgent idea: adaptive computation time [4, 17], universal transformers [2], recurrent networks that extrapolate by iterating longer [18, 19], looped transformers [20], recurrent-depth language models [3], and tiny recursive reasoners [21, 22], evaluated on algorithmic suites such as CLRS [23]. Claims in this literature share a load-bearing assumption: that the iterated module actually computes. Our protocol (Section 7) is a cheap audit of exactly that assumption, and our case study shows the assumption can fail silently while every surface metric looks healthy.

Negative results and diagnostics. Our probe methodology follows the linear-probing tradition [24] with the caution urged by control tasks [25] — our blow-up regime is a concrete instance of a probe finding “signal” that is not computation. The mechanism we identify is a module-level relative of gradient starvation [26] and of shortcut learning [27]: a direct pathway that suffices for the loss removes the incentive for the expensive pathway to learn. We use the term in this architectural sense rather than the feature-dynamics sense of Pezeshki et al. [26].

3 Model and Tasks

3.1 Port-Hamiltonian anchored DEQ

The implicit state is phase-space-structured, $z = (q, p)$, evolved by a damped, anchored, symplectic-Euler-style update conditioned on a context vector c :

$$q_{t+1} = q_t + \epsilon \nabla_p K(p_t, c) - \epsilon r(c) p_t + \epsilon r_a (q_0 - q_t), \quad (1)$$

$$p_{t+1} = p_t - \epsilon \nabla_q U(q_{t+1}, c) - \epsilon r(c) q_{t+1} + \epsilon r_a (p_0 - p_t), \quad (2)$$

where K and U are learned scalar (kinetic and potential) MLPs, $r(c) = \text{softplus}(w_r^\top c + b_r)$ is a learned context-conditioned damping coefficient, $\epsilon = 0.05$, and r_a is the anchor strength (0.02 by default; 0 in arm C) pulling the state toward the learned initialization $z_0 = (q_0, p_0)$. Weight matrices of K and U are spectrally clamped to norm ≤ 0.95 after every update. The fixed point of Eqs. (1)–(2) is solved by Anderson acceleration ($m=5$, cap 300 iterations) from a start point that equals z_0 in the standard design, to a per-example adaptive tolerance $10^{-(5\sigma(w_t^\top c + b_t) + 1)} \in [10^{-6}, 10^{-1}]$. Gradients pass through the implicit function theorem via an adjoint Anderson solve. The initialization network maps context to phase space: q_0 by a linear read-out of the premise/graph context, and p_0 through a goal-gated channel in which a query-key gate modulates value projections of the context; we study both a scalar-gate and an elementwise-gate variant of this channel (Section 8). The decoder is a two-layer GELU MLP reading $[h^*; c_{\text{ctx}}; c_{\text{query}}]$. Backward-pass coverage of z_0 differs between our two pipelines and matters for the mechanism; see Section 5.3.

3.2 Tasks and positive controls

ProofWriter entailment. Binary entailment on ProofWriter [28], stratified over proof depths $\{1, 2, 3, 5\}$ (800 per depth where available; depth-5 capped at 697), yielding 2477 train / 620 test examples with label balance 0.495/0.477. Premises and goals are embedded by frozen DeBERTa-v3-base [29]; embeddings are cached, so every method reads identical inputs. **Graph reachability.** Synthetic k -hop reachability: random directed graphs on 16 nodes with out-degree 2; a query (s, t, k) asks whether t is reachable from s within $k \leq 8$ hops, labeled by breadth-first search and balanced exactly (6,000 train / 1,000 test; 30,000 train in arm E). Every instance carries its true BFS distance,




Figure 1: Training-time displacement between the solved equilibrium and the solver’s start point, graph task, all seeds. Arms A (coupled) and C (anchor removed) sit at zero; arm B (retrained from noisy starts) sits exactly on the injected-noise norm. The equilibrium is the initialization.

giving a ground-truth difficulty signal for the iteration analysis (Section 4.3). **Positive controls.** A two-layer MLP on the identical cached inputs is trained fresh per seed alongside every pipeline run ($90.39 \pm 0.33\%$ on ProofWriter; the analogous context-only model is arm D on graphs), establishing both that the data pipeline supports learning and what the implicit apparatus must beat.¹

4 Results: Two Degenerate Regimes

4.1 The equilibrium is the start point

Figure 1 shows the training-time displacement $\|h^* - \text{start}\|$ for every seed of three controlled arms on the graph task. In arm A (the original design: solver started at the anchor z_0), displacement is numerically zero throughout training and remains so at test time (mean 6.9×10^{-7} , pooled over 5 seeds). Removing the anchor entirely (arm C, $r_{\text{anchor}} = 0$) changes nothing: per-seed accuracies are identical to arm A and loss curves agree to four decimal places, with displacement differing only at the seventh decimal. The anchoring term is causally irrelevant to the collapse. Most tellingly, when the solver is *retrained* from noisy starts $z_0 + \sigma\epsilon$ (arm B, $\sigma=2$, ϵ resampled every forward pass), training displacement sits throughout on the theoretical norm of the injected noise, $\sigma \mathbb{E} \|\epsilon\| \approx 7.87$: the solved “equilibrium” is the start point, wherever the start is placed. The same invariance holds across a five-value sweep of the anchor strength $r_{\text{anchor}} \in \{0.005, 0.02, 0.05, 0.1, 0.2\}$ (displacement 2×10^{-6} at every value) and persists when the displacement penalty is removed from the loss, ruling out the regularizer as the cause. On ProofWriter, 4 of 5 seeds show the same signature (displacement ≤ 0.12 ; Table 2); the fifth exhibits the second regime, described in Section 4.4.

¹Working notes from an unpreserved earlier session recorded an 81.72% baseline; it failed re-verification against executed outputs and is superseded by the fresh per-seed values reported here (see Limitations, reproducibility discipline).

Table 1: Graph reachability, mean $\pm$ std over seeds. “Start-sub” bypasses the solver, feeding its start point to the decoder. The implicit computation contributes +0.00pp everywhere; the context-only arm matches the full pipeline. All arms reach 100% train accuracy.

Arm	Full	Start-sub	Gap (pp)	h^* -zeroed	$\ h^* - z_0\ $
A coupled (5s)	61.80 ± 1.34	61.80 ± 1.34	0.00 ± 0.00	60.92 ± 1.01	7×10^{-7}
B decoupled retrain (5s)	61.82 ± 0.96	61.82 ± 0.96	0.00 ± 0.00	61.84 ± 0.96	3×10^{-7}
C no anchor (5s)	61.80 ± 1.34	61.80 ± 1.34	0.00 ± 0.00	60.92 ± 1.01	7×10^{-7}
E 30k graphs (3s)	64.20 ± 0.79	64.20 ± 0.79	0.00 ± 0.00	64.37 ± 1.32	7×10^{-7}
D context-only (5s)	61.52 ± 0.72	—	—	—	—

Table 2: ProofWriter, per seed (pooling would average two distinct regimes). Seeds 1, 2, 4, 5: lazy identity collapse. Seed 3: solver blow-up — start-substitution falls below chance while *zeroing* h^* improves over the full model. The fresh per-seed MLP baseline (right) is never beaten.

Seed	Full	Start-sub	Gap (pp)	h^* -zeroed	$\ h^* - z_0\ $	MLP baseline
1	91.77	91.77	+0.00	47.74	2×10^{-6}	90.65
2	91.29	91.61	−0.32	53.39	0.12	90.32
3	81.77	32.58	+49.19	83.23	171.43	90.00
4	90.32	90.32	+0.00	47.74	2×10^{-6}	90.81
5	91.77	91.77	+0.00	52.26	1×10^{-7}	90.16

4.2 Bypassing the solver changes nothing

The unconfounded measure of the DEQ’s contribution is the *substitution test*: feed the solver’s start point directly to the decoder, skipping the fixed-point solve entirely. Table 1 reports the result on the graph task: the gap between the full pipeline and the solver-bypassed pipeline is +0.00pp in every arm and every seed — including arm B, which was explicitly retrained to break the start–anchor coupling, and arm E, which scales training data $5\times$. A context-only baseline (arm D: identical embedding and decoder, no initializer, no DEQ) matches the full pipeline (61.52 ± 0.72 vs. 61.80 ± 1.34), so the *entire* implicit apparatus is dispensable. On ProofWriter the substitution gap is ≤ 0.32 pp in the 4 collapsed seeds (Table 2).

4.3 Iteration counts ignore difficulty

If the implicit layer performed input-adaptive computation, harder instances should require more solver iterations. They do not (Figure 3). On the graph task — where every instance carries a BFS-verified ground-truth distance — the correlation between iteration count K and true distance is 0.009 (arm A, 5 seeds pooled; −0.032 at $5\times$ data), with 99.9% of solves running to the iteration cap regardless of difficulty. Accuracy is likewise constant across solver tolerances from 10^{-2} to 10^{-6} : since the state never moves, solve precision is irrelevant.

4.4 The minority regime: solver blow-up

One ProofWriter seed in five escapes the identity collapse — by diverging. Seed 3 ends training with mean displacement 171.4; its start-substituted accuracy falls to 32.58%, *below* the 47.7% majority-class floor, while zeroing h^* *raises* accuracy from 81.77% to 83.23% (Table 2, Figure 4). The interpretation is unflattering to the implicit layer in a different way: the non-converged output




Figure 2: Substitution and ablation accuracies. Left: graph task (mean±std); the dashed line is the context-only arm D. Right: ProofWriter per seed; seed 3 is the blow-up regime.

acts as a high-variance channel the decoder has co-adapted to, and removing it outright helps. We observed the same instability earlier when an elementwise initializer gate increased gradient flow by an order of magnitude: solver residuals grew from 0.09 to a mean of 29.1 (max 48.3) while a probe on h^* appeared to find signal — signal riding on divergence, not equilibrium (Appendix B). The two regimes jointly close the escape routes: the equilibrium either stays at its start (and contributes nothing) or leaves it by diverging (and contributes noise).

4.5 The apparatus never beats a two-layer MLP

Finally, the positive controls double as an upper-bound check. On ProofWriter, a two-layer MLP on the same frozen embeddings reaches $90.39 \pm 0.33\%$ (fresh initialization per seed); the full physics-structured pipeline reaches $89.39 \pm 4.30\%$. On the graph task, the context-only arm matches the full pipeline. Across both tasks, at no point does the port-Hamiltonian DEQ apparatus outperform the simplest baseline it was meant to improve upon.

5 Mechanism

5.1 Where the signal actually flows

On ProofWriter, all task-relevant information travels through the feed-forward initializer’s goal-gated momentum channel and the decoder’s direct view of the context. Zeroing the p_0 (goal-gated) half of the initializer output at inference drops accuracy from 91.77% to 47.74% — the majority-class floor — while zeroing the q_0 half changes nothing (91.77%). Accuracy from the start point alone is essentially flat in proof depth (92.5/92.5/91.9/90.0% at depths 1/2/3/5), consistent with a task solvable by a shallow feed-forward read-out rather than iterative inference. A fresh classifier probing the frozen equilibrium h^* alone converges to the label-entropy floor (loss $\approx \ln 2$; test accuracy 51.45%); the equilibrium carries no task signal beyond what it copies from its start.




Figure 3: Left: iteration histograms (arm B’s bimodality reflects noisy starts that occasionally satisfy the tolerance immediately). Right: mean K vs. ground-truth BFS distance; an adaptive solver would slope upward.

5.2 Causal dissection: what does and does not cause the collapse

The design invites a seductive mechanistic story: the anchor both initializes the solver and attracts it, so the solver starts where it is pulled back to, displacement is identically zero, and the Hamiltonian networks never receive gradients away from the anchor — a self-fulfilling no-op. Two controlled arms falsify the causal part of this story. Removing the anchor entirely (arm C) reproduces arm A’s per-seed results to reported precision. Retraining with the start decoupled from the anchor (arm B) — the intervention the self-fulfilling-no-op account predicts should rescue the dynamics — does not: the solver remains a near-identity map on its (now noisy) start, and the decoder’s adaptation is to *ignore* the h^* channel, as witnessed by its zeroing ablation matching full accuracy exactly (61.84 ± 0.96 vs. 61.82 ± 0.96). Forcing displaced starts at inference on the trained coupled model completes the picture: displacement becomes genuine (≈ 7.9) but accuracy *drops* (56.06 ± 1.26 vs. 61.80 ± 1.34), so the off-start regions of the learned vector field contain no structure — capability absent, not suppressed.

5.3 Gradient starvation: two routes to the same no-op

The deeper cause is that nothing in the gradient structure ever rewards genuine iterative computation, and a near-identity fixed-point map is a stable optimum. Both tasks are solvable from the direct context pathway (Section 4.5), so the loss can be driven to zero while the equilibrium contributes nothing; once it is, the gradient reaching the dynamics networks through the implicit function theorem is proportional to a decoder sensitivity that the decoder itself is free to remove. The two pipelines realize this starvation by different routes, which is worth stating precisely because it shows the phenomenon is not an implementation quirk. In the ProofWriter pipeline, the backward pass extends to z_0 , the initializer trains, and z_0 genuinely carries signal (Section 5.1) — but the solve on top of it is an identity. In the graph pipeline, the custom VJP returns zero cotangents for both z_0 and the context, so the initializer is *gradient-isolated* and remains at its random initialization; the decoder learns from its direct context inputs, and the equilibrium faithfully copies an untrained




Figure 4: ProofWriter seeds in the (displacement, substitution-gap) plane. Four seeds collapse to the start; one diverges. Neither regime performs useful implicit computation.

random projection. Same symptom, two gradient-starvation routes: one starves the dynamics of incentive, the other additionally starves the start of training. At initialization the port-Hamiltonian update is already near-identity (small MLP outputs, step size $\epsilon = 0.05$, spectral norms clamped below 0.95), so “stay put” is where training begins — and nothing ever pays it to leave, except, occasionally, instability (Section 4.4).

6 A Note on Confounded Ablations

Our own first attempt at measuring the DEQ’s contribution was wrong in an instructive way. The natural ablation — zero the equilibrium h^* at the decoder input — produced wildly seed-dependent deltas on ProofWriter (+44.0, +19.5, +0.0, +44.0, +39.5 pp across five seeds), which superficially read as “the DEQ sometimes matters a great deal.” It does not. Zeroing an internal representation removes *two* things at once: the computation performed on the representation, and the input pathway flowing *through* it — here, the start point z_0 that the identity-solve copies to the decoder. The seed variance in the zeroing deltas tracks how each decoder happens to distribute weight between its h^* channel and its direct context channels, not what the solver computes. The substitution test isolates the computation by keeping the pathway and bypassing only the solve, and it returns 0.00 pp with negligible variance in every collapsed run. The blow-up seed makes the confound vivid from the other side: there, zeroing *helps* (Section 4.4), because the “pathway” being removed is co-adapted noise. We suggest the general rule: whenever an iterative or implicit module is credited with computation, the control should substitute the module’s input for its output — not delete the channel.

7 A Diagnostic Protocol for Implicit Computation

The failure we document is silent: training converges, accuracy is respectable, solver telemetry looks plausible, and the model card would read “port-Hamiltonian deep equilibrium reasoner.” Nothing in the standard training signals reveals that the equilibrium computation is a no-op. We

therefore distill our analysis into four inexpensive tests, applicable to any architecture that claims iterative or implicit computation over a learned initialization — DEQs, looped transformers, and recurrent-depth models alike. On our tasks the full protocol runs in minutes on a free Colab GPU.

T1 — Substitution (the primary test). At evaluation, replace the module’s output with its input: feed the solver’s start point (or the loop’s initial state) directly to the decoder. *Fail* if accuracy is unchanged (the module computes nothing); in our runs the gap was +0.00 pp across 18 of 19 training runs spanning two tasks. This test — not the zeroing ablation — isolates the computation from the pathway (Section 6).

T2 — Representation probe against the entropy floor. Train a fresh classifier on the module’s output *alone*, frozen. *Fail* if probe loss converges to the label-entropy floor ($\ln 2$ for balanced binary tasks; 51.45% accuracy in our case). A passing probe with a *diverging* solver should be treated as a false positive: verify residuals before crediting the representation (our gate-fix run probed at 88.9% while mean residuals sat at 29.1).

T3 — Iteration variance against difficulty. On a task with a controllable or verifiable difficulty measure (here, ground-truth BFS distance), correlate solver iterations with difficulty. *Fail* if the correlation is ≈ 0 or iteration counts are degenerate (pinned at the cap, or at immediate exit): the model is not allocating computation to instances. Our runs: $\text{corr} = 0.009$ with 99.9% of solves at the cap.

T4 — Forced displacement. Start the solver away from its learned initialization at inference (we used Gaussian offsets with σ matched to the state scale). *Fail-absent* if accuracy degrades with genuine displacement — the off-initialization vector field was never trained and contains no structure — as opposed to *fail-suppressed*, where accuracy is maintained, indicating a capability the training regime masks. Our runs degrade (61.8% $\rightarrow$ 56.1%), and the retraining version of this test (arm B) shows the decoder simply learns to ignore the module.

A model that passes all four tests may still not “reason,” but a model that fails T1 certainly does not, whatever its architecture diagram promises. We suggest T1 as a standard reported control whenever implicit or recurrent computation is part of a paper’s claimed contribution.

8 Secondary Finding: Expressive Gates Destabilize the Solve

During the ProofWriter study we replaced the scalar goal gate in the initialization network with an elementwise gate, which increased the gradient norm reaching the initializer by an order of magnitude (0.0027 $\rightarrow$ 0.026) — and destabilized the solve: mean final residuals grew from 0.09 to 29.1 (max 48.3), so the returned h^* was no longer an equilibrium at all. Instructively, a probe on this non-converged h^* reached 88.9%: apparent representational “signal” riding entirely on divergence, foreshadowing the seed-3 blow-up regime of Section 4.4. An instrumented run captured the onset window (epochs 74–90): peak state magnitude grew 6.9 $\rightarrow$ 21.7, initializer gradients spiked to 1.7, and accuracy fell from 92% to 62% before partially recovering (Appendix B). The general lesson is a trade-off worth naming: interventions that increase gradient flow into the initialization of an implicit model can buy trainability at the price of solver validity — and standard telemetry (accuracy, probe scores) does not distinguish the two.

9 Limitations

Scope of the negative claim. Our evidence concerns one architecture family (port-Hamiltonian dynamics with a learned initializer, solved by Anderson acceleration with $m=5$ and a 300-iteration cap) on two binary classification tasks. We claim that *this* design performs no useful implicit computation and that the mechanism is gradient starvation; we do not claim that all DEQs, looped models, or physics-structured implicit layers must fail. Indeed, the diagnostic framing exists precisely because the failure is silent and design-specific. Tasks where the direct context pathway *cannot* solve the problem — unlike both of ours (Section 4.5) — may exert the gradient pressure our setup lacks; establishing this is future work.

The two pipelines differ in backward-pass coverage. In the ProofWriter pipeline the backward pass extends to the initializer; in the graph pipeline the custom VJP returns zero cotangents for z_0 and the context, leaving the initializer untrained (Section 5.3). We report this as two routes to the same starvation, but it means the fully-differentiated configuration was tested on only one of the two tasks. The substitution result is unaffected — it is a forward-pass property and holds identically in both pipelines.

Generalization on the synthetic task. All graph-task arms memorize (100% train accuracy) and generalize weakly ($\sim 62\%$ at 6,000 graphs, 64.2% at 30,000). The lazy-identity result is orthogonal to this gap — the substitution test compares two decoders reading the same representations — but task-level accuracies on the graph benchmark should not be read as statements about reachability reasoning in general.

The blow-up regime is characterized from few events. The divergent regime appears in 1 of 5 ProofWriter seeds (plus, mechanistically consistent, in our earlier gate-modified run). We can assert its existence and its signature — below-chance substitution, zeroing that improves accuracy, displacement orders of magnitude above the collapsed regime — but not its frequency as a function of hyperparameters.

Frozen encoders. Both tasks use frozen or from-scratch context encoders (DeBERTa-v3-base embeddings; a small trained embedding for graphs). End-to-end finetuning of a large encoder could change the gradient landscape reaching the implicit layer.

Reproducibility discipline. All numbers in this paper are recomputed from raw per-run logs by a released analysis script; two figures quoted in our own working notes from an unpreserved earlier session (a 81.72% baseline; a 51.29% probe) were retired when they failed re-verification, and the fresh multi-seed values are reported instead. Runs cost 2–6 minutes each on a free Colab GPU; optimizer and solver settings are identical across pipelines (AdamW, 3×10^{-4} , weight decay 10^{-4} , global-norm clip 1.0).

10 Conclusion

We set out to build a physics-structured implicit reasoner and instead built a very careful copy machine. The equilibrium of our port-Hamiltonian DEQ is its initialization — under the anchored design, without the anchor, and even when retrained from randomized starts — and the one training run that escaped this regime did so by diverging into decoder-adapted noise. Nothing in ordinary

training telemetry reveals any of this; it took substitution tests, gradient-path inspection, and controlled retraining arms to see it.

We draw three implications. First, for the growing literature that credits looped, recurrent-depth, and implicit models with *reasoning*: the credit assumes the iterated module computes, and that assumption is cheap to test and capable of failing silently — we suggest the substitution test (T1) as a standard reported control. Second, for architecture design: when a task is solvable through a direct pathway, adding an expensive iterative module does not merely risk wasted parameters; the optimizer will actively route around it, and stabilizing structure can make the do-nothing solution easier to reach. Gradient incentives, not architectural affordances, determine whether iteration is used. Third, for practice: negative results of this shape are inexpensive to establish rigorously — our entire evidence base cost minutes per run on a free GPU — and the diagnostic protocol transfers directly to models far larger than ours.

Reproducibility. All experiments run on a single free Colab GPU (per-run cost: $\sim$ 2–6 minutes training). Code, configs, seeds, and raw logs: <https://github.com/joyjeet-singh/lazy-identity-deq>

References

- [1] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, 2019.
- [2] Mostafa Dehghani, Stephan Gouws, Oriol Vinyals, Jakob Uszkoreit, and Łukasz Kaiser. Universal transformers. In *International Conference on Learning Representations*, 2019.
- [3] Jonas Geiping, Sean McLeish, Neel Jain, John Kirchenbauer, Siddharth Singh, Bhavya Kailkhura, Abhinav Bhatele, and Tom Goldstein. Scaling up test-time compute with latent reasoning: A recurrent depth approach. *arXiv preprint arXiv:2502.05171*, 2025.
- [4] Alex Graves. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*, 2016.
- [5] Samuel Greydanus, Misko Dzamba, and Jason Yosinski. Hamiltonian neural networks. In *Advances in Neural Information Processing Systems*, 2019.
- [6] Pengzhan Jin, Zhen Zhang, Aiqing Zhu, Yifa Tang, and George Em Karniadakis. SympNets: Intrinsic structure-preserving symplectic networks for identifying Hamiltonian systems. *Neural Networks*, 132:166–179, 2020.
- [7] Shaan A. Desai, Marios Mattheakis, David Sondak, Pavlos Protopapas, and Stephen J. Roberts. Port-Hamiltonian neural networks for learning explicit time-dependent dynamical systems. *Physical Review E*, 104(3):034312, 2021.
- [8] Shaojie Bai, Vladlen Koltun, and J. Zico Kolter. Multiscale deep equilibrium models. In *Advances in Neural Information Processing Systems*, 2020.
- [9] Laurent El Ghaoui, Fangda Gu, Bertrand Travacca, Armin Askari, and Alicia Tsai. Implicit deep learning. *SIAM Journal on Mathematics of Data Science*, 3(3):930–958, 2021.
- [10] Ezra Winston and J. Zico Kolter. Monotone operator equilibrium networks. In *Advances in Neural Information Processing Systems*, 2020.

- [11] Donald G. Anderson. Iterative procedures for nonlinear integral equations. *Journal of the ACM*, 12(4):547–560, 1965.
- [12] Shaojie Bai, Vladlen Koltun, and J. Zico Kolter. Stabilizing equilibrium models by Jacobian regularization. In *International Conference on Machine Learning*, 2021.
- [13] Samy Wu Fung, Howard Heaton, Qiuwei Li, Daniel McKenzie, Stanley Osher, and Wotao Yin. JFB: Jacobian-free backpropagation for implicit networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2022. arXiv:2103.12803.
- [14] Cem Anil, Ashwini Pople, Kaiqu Liang, Johannes Treutlein, Yuhuai Wu, Shaojie Bai, J. Zico Kolter, and Roger B. Grosse. Path independent equilibrium models can better exploit test-time computation. In *Advances in Neural Information Processing Systems*, 2022.
- [15] Miles Cranmer, Sam Greydanus, Stephan Hoyer, Peter Battaglia, David Spergel, and Shirley Ho. Lagrangian neural networks. *arXiv preprint arXiv:2003.04630*, 2020.
- [16] Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David Duvenaud. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems*, 2018. arXiv:1806.07366.
- [17] Andrea Banino, Jan Balaguer, and Charles Blundell. PonderNet: Learning to ponder. *arXiv preprint arXiv:2107.05407*, 2021.
- [18] Avi Schwarzschild, Eitan Borgnia, Arjun Gupta, Furong Huang, Uzi Vishkin, Micah Goldblum, and Tom Goldstein. Can you learn an algorithm? Generalizing from easy to hard problems with recurrent networks. In *Advances in Neural Information Processing Systems*, 2021.
- [19] Arpit Bansal, Avi Schwarzschild, Eitan Borgnia, Zeyad Emam, Furong Huang, Micah Goldblum, and Tom Goldstein. End-to-end algorithm synthesis with recurrent networks: Extrapolation without overthinking. In *Advances in Neural Information Processing Systems*, 2022.
- [20] Liu Yang, Kangwook Lee, Robert D. Nowak, and Dimitris Papailiopoulos. Looped transformers are better at learning learning algorithms. In *International Conference on Learning Representations*, 2024.
- [21] Guan Wang, Jin Li, Yuhao Sun, Xing Chen, Changling Liu, Yue Wu, Meng Lu, Sen Song, and Yasin Abbasi Yadkori. Hierarchical reasoning model. *arXiv preprint arXiv:2506.21734*, 2025.
- [22] Alexia Jolicoeur-Martineau. Less is more: Recursive reasoning with tiny networks. *arXiv preprint arXiv:2510.04871*, 2025.
- [23] Petar Veličković, Adrià Puigdomènech Badia, David Budden, Razvan Pascanu, Andrea Banino, Misha Dashevskiy, Raia Hadsell, and Charles Blundell. The CLRS algorithmic reasoning benchmark. In *International Conference on Machine Learning*, 2022.
- [24] Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. *arXiv preprint arXiv:1610.01644*, 2016. v1 October 2016.
- [25] John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of EMNLP-IJCNLP*, 2019.

- [26] Mohammad Pezeshki, Sékou-Oumar Kaba, Yoshua Bengio, Aaron Courville, Doina Precup, and Guillaume Lajoie. Gradient starvation: A learning proclivity in neural networks. In *Advances in Neural Information Processing Systems*, 2021.
- [27] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020. arXiv:2004.07780.
- [28] Oyvind Tafjord, Bhavana Dalvi Mishra, and Peter Clark. ProofWriter: Generating implications, proofs, and abductive statements over natural language. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP*, 2021.
- [29] Pengcheng He, Jianfeng Gao, and Weizhu Chen. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. In *International Conference on Learning Representations*, 2023.

A Why Physics Structure Seemed Worth Trying

The port-Hamiltonian parameterization was not chosen arbitrarily. In a standalone rollout experiment on learned two-body-style dynamics, we compared long-horizon energy behavior of an unconstrained baseline network against the structure–dissipation parameterization used in this paper. Over a 5,000-step rollout the baseline’s normalized Hamiltonian drift saturates at 6.2×10^{25} (3.5×10^{17} already by step 501), while the port-Hamiltonian network holds 2.3×10^3 (and 2.3×10^{-2} at step 501) — roughly twenty orders of magnitude of improvement in energy control. This is the standard, genuine benefit of the inductive bias, and it is precisely why the main text’s finding is worth reporting: stability of the dynamics and usefulness of the equilibrium are different properties, and the first does not purchase the second.

B Training Curves and Instability Window

Figure 5 shows training loss and accuracy for every seed of every graph-task arm; all runs train to convergence (100% train accuracy), closing the under-training explanation. For the ProofWriter gate-variant instability (Section 8), the instrumented onset window recorded per-epoch peak state magnitude $\max |h^*|$ rising $6.9 \rightarrow 21.7$ across epochs 74–90 with initializer gradient spikes to 1.7 and a $92\% \rightarrow 62\%$ accuracy excursion; solver mean residuals in this variant reached 29.1 (cf. 0.09 in the standard pipeline), i.e. returned states were not equilibria.

C Full Experimental Details

Optimization. AdamW, learning rate 3×10^{-4} , weight decay 10^{-4} , global-norm clip 1.0; batch size 64; 100 epochs (ProofWriter) / 80 epochs (graphs); ProofWriter uses an auxiliary-weight ramp $\alpha = 0.1 \min(1, \text{epoch}/20)$. **Solver.** Anderson acceleration, memory $m=5$, iteration cap 300, Tikhonov-regularized mixing with coefficient clipping to $[-3, 3]$ and state clipping to $[-50, 50]$; adjoint solve: cap 100, tolerance 10^{-4} . **Model sizes.** Phase dimension 6 (ProofWriter) / 8 (graphs); context dimension 128; kinetic and potential MLPs with hidden width 32 and softplus activations; decoder hidden width 64 (GELU). **Data.** ProofWriter examples scanned up to a 60,000-row cap and bucketed by proof depth to the splits in Section 3; graph datasets generated with a fixed data




Figure 5: Training curves, graph task, all arms and seeds.

seed (42), exactly class-balanced, with dataset SHA-256 recorded in every run manifest. **Hardware and cost.** Single free-tier Colab GPU; ~ 2 minutes per ProofWriter run, ~ 5 – 6 minutes per graph run (~ 26 at 30k examples). **Logging.** Every run writes a manifest (library versions, device, config hash, dataset hash) and per-epoch JSONL records; all tables and figures in this paper are generated from those files by the released analysis scripts.